

01d - Introduction - OpenML

January 18, 2017

1 OpenML in Python

OpenML is an online collaboration platform for machine learning:

- Share/reuse machine learning datasets, algorithms, models, experiments
- Well documented/annotated datasets, uniform access
- APIs in Java, R, Python*,... to download/upload everything
- Better reproducibility of experiments, reuse of machine learning models
- Works well with machine learning libraries such as scikit-learn
- Large scale benchmarking, compare to state of the art

```
In [24]: %matplotlib inline
         from preamble import * # Ignore, this is just to make code cleaner
         InteractiveShell.ast_node_interactivity = "all"
         HTML(''<style>html, body{overflow: visible !important} .CodeMirror{min-width:105% !imp
```

```
Out[24]: <IPython.core.display.HTML object>
```

1.1 Authentication

- Create an OpenML account (free) on <http://www.openml.org>.
- After logging in, open your account page (avatar on the top right)
- Open 'Account Settings', then 'API authentication' to find your API key.

There are two ways to authenticate:

- Create a plain text file `~/.openml/config` with the line `'apikey=MYKEY'`, replacing MYKEY with your API key.
- Run the code below, replacing 'MYKEY' with your API key.

```
In [25]: # Uncomment and run this to authenticate. Don't share your API key!
         # oml.config.apikey = os.environ.get('OPENMLKEY', 'MYKEY')
```

2 Data sets

We can list, select, and download all OpenML datasets

2.0.1 List datasets

```
In [26]: datalist = oml.datasets.list_datasets() # Returns a dict
         datalist = pd.DataFrame.from_dict(datalist, orient='index') # Create a DataFrame
         print("First 10 of %s datasets..." % len(datalist))
         datalist[:10][['did', 'name', 'NumberOfInstances',
                        'NumberOfFeatures', 'NumberOfClasses']]
```

First 10 of 19492 datasets...

```
Out [26]:
```

	did	name	NumberOfInstances	NumberOfFeatures	NumberOfClasses
1	1	anneal	898	39	6
2	2	anneal	898	39	6
3	3	kr-vs-kp	3196	37	2
4	4	labor	57	17	2
5	5	arrhythmia	452	280	16
6	6	letter	20000	17	26
7	7	audiology	226	70	24
8	8	liver-disorders	345	7	-1
9	9	autos	205	26	7
10	10	lymph	148	19	4

There are many properties that we can query

```
In [27]: list(datalist)
         datalist = datalist[['did', 'name', 'NumberOfInstances',
                              'NumberOfFeatures', 'NumberOfClasses']]
```

```
Out [27]: ['status',
            'NumberOfSymbolicFeatures',
            'did',
            'NumberOfInstances',
            'NumberOfFeatures',
            'MinorityClassSize',
            'NumberOfNumericFeatures',
            'MajorityClassSize',
            'name',
            'NumberOfMissingValues',
            'format',
            'NumberOfInstancesWithMissingValues',
            'NumberOfClasses',
            'MaxNominalAttDistinctValues']
```

and we can filter or sort on all of them

```
In [28]: datalist[datalist.NumberOfInstances>10000
               ].sort(['NumberOfInstances'])[:20]
```

```
Out[28]:
```

	did	name	NumberOfInstances	\
23515	23515	sulfur	10081	
372	372	internet_usage	10108	
981	981	kdd_internet_usage	10108	
1536	1536	volcanoes-b6	10130	
4562	4562	InternetUsage	10168	
1531	1531	volcanoes-b1	10176	
1534	1534	volcanoes-b4	10190	
1459	1459	artificial-characters	10218	
1478	1478	har	10299	
1533	1533	volcanoes-b3	10386	
1532	1532	volcanoes-b2	10668	
1053	1053	jm1	10885	
1414	1414	Kaggle_bike_sharing_demand_challenge	10886	
1044	1044	eye_movements	10936	
1019	1019	pendigits	10992	
32	32	pendigits	10992	
4534	4534	PhishingWebsites	11055	
399	399	ohscal.wc	11162	
310	310	mammography	11183	
1568	1568	nursery	12958	

	NumberOfFeatures	NumberOfClasses
23515	7	-1
372	72	46
981	69	2
1536	4	5
4562	72	-1
1531	4	5
1534	4	5
1459	8	10
1478	562	6
1533	4	5
1532	4	5
1053	22	2
1414	12	-1
1044	28	3
1019	17	2
32	17	10
4534	31	2
399	11466	10
310	7	2
1568	9	4

or find specific ones

```
In [29]: datalist.query('name == "eeg-eye-state"')
```

```
Out[29]:
```

did	name	NumberOfInstances	NumberOfFeatures	\
-----	------	-------------------	------------------	---

1471	1471	eeg-eye-state	14980	15
		NumberOfClasses		
1471		2		

```
In [30]: datalist.query('NumberOfClasses > 50')
```

```
Out[30]:
```

	did	name	NumberOfInstances	NumberOfFeatures	\
1491	1491	one-hundred-plants-margin	1600	65	
1492	1492	one-hundred-plants-shape	1600	65	
1493	1493	one-hundred-plants-texture	1599	65	
4546	4546	Plants	44940	16	
4552	4552	BachChoralHarmony	5665	17	
		NumberOfClasses			
1491		100			
1492		100			
1493		100			
4546		57			
4552		102			

Download a specific dataset. This is done based on the dataset ID (called 'did').

```
In [31]: dataset = oml.datasets.get_dataset(1471)
```

```
print("This is dataset '%s', the target feature is '%s'" %
      (dataset.name, dataset.default_target_attribute))
print("URL: %s" % dataset.url)
print(dataset.description[:500])
```

```
This is dataset 'eeg-eye-state', the target feature is 'Class'
URL: http://www.openml.org/data/download/1587924/php1E7q6h
**Author**: Oliver Roesler, it12148 '@' lehre.dhbw-stuttgart.de
**Source**: [UCI] (https://archive.ics.uci.edu/ml/datasets/EEG+Eye+State), Baden-Wuerttemberg, Co
**Please cite**:
```

All data is from one continuous EEG measurement with the Emotiv EEG Neuroheadset. The duration o

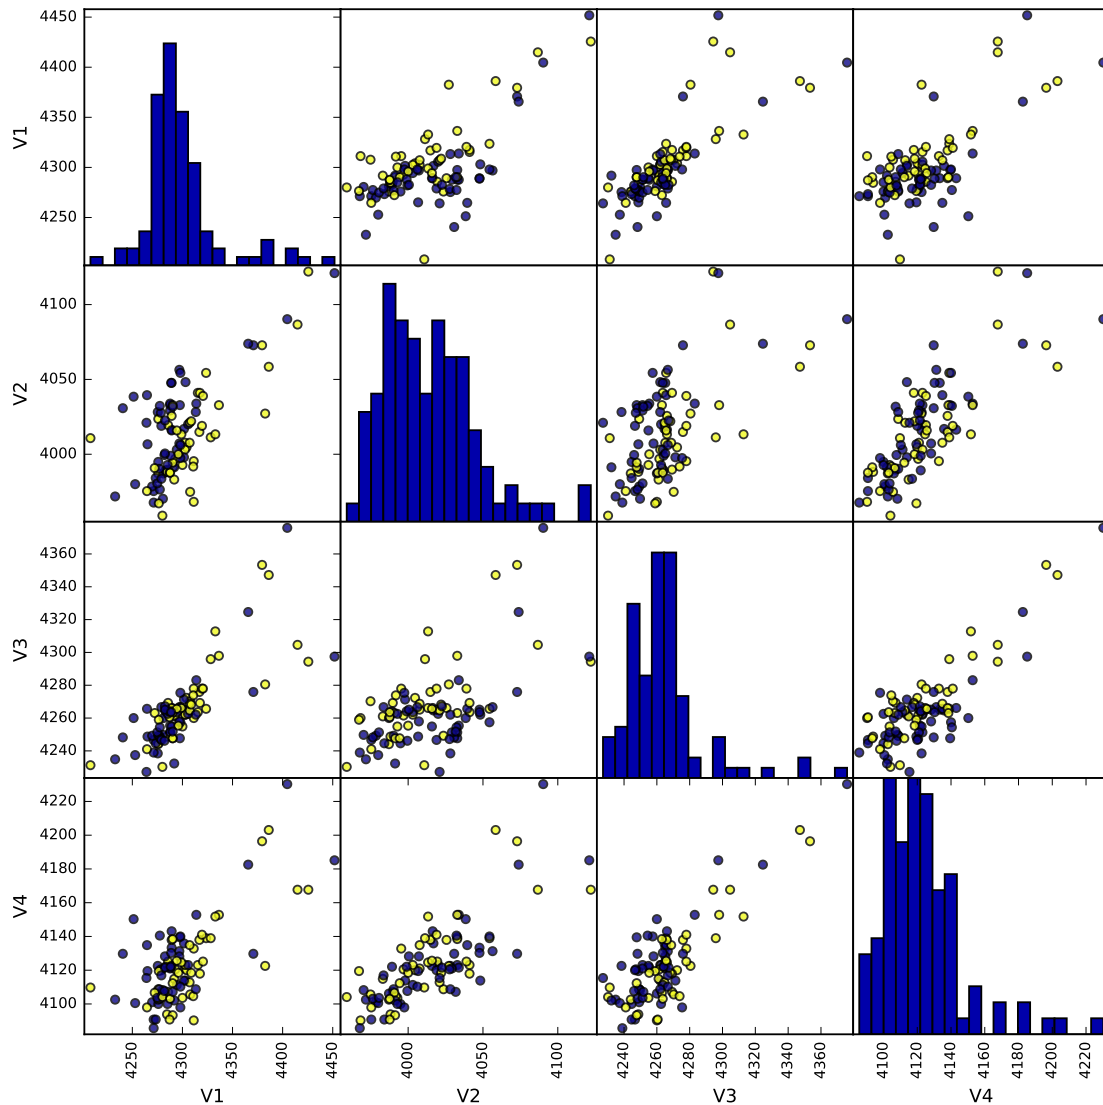
Convert the data to a DataFrame for easier processing/plotting

```
In [32]: X, y, attribute_names = dataset.get_data(
          target=dataset.default_target_attribute,
          return_attribute_names=True)
eeg = pd.DataFrame(X, columns=attribute_names)
eeg['class'] = y
print(eeg[:10])
```

	V1	V2	V3	V4	...	V12	V13	V14	class
0	4329.23	4009.23	4289.23	4148.21	...	4280.51	4635.90	4393.85	0
1	4324.62	4004.62	4293.85	4148.72	...	4279.49	4632.82	4384.10	0
2	4327.69	4006.67	4295.38	4156.41	...	4282.05	4628.72	4389.23	0
3	4328.72	4011.79	4296.41	4155.90	...	4287.69	4632.31	4396.41	0
4	4326.15	4011.79	4292.31	4151.28	...	4288.21	4632.82	4398.46	0
5	4321.03	4004.62	4284.10	4153.33	...	4281.03	4628.21	4389.74	0
6	4319.49	4001.03	4280.51	4151.79	...	4269.74	4625.13	4378.46	0
7	4325.64	4006.67	4278.46	4143.08	...	4266.67	4622.05	4380.51	0
8	4326.15	4010.77	4276.41	4139.49	...	4273.85	4627.18	4389.74	0
9	4326.15	4011.28	4276.92	4142.05	...	4277.95	4637.44	4393.33	0

[10 rows x 15 columns]

```
In [43]: eegs = eeg.sample(n=1000)
         _ = pd.scatter_matrix(eegs.iloc[:100,:4], c=eegs[:100]['class'], figsize=(10, 10),
                               marker='o', hist_kwds={'bins': 20},
                               alpha=.8, cmap='plasma')
```



2.1 Train models

Train a scikit-learn model on the data manually

In [34]: `from sklearn import neighbors`

```
dataset = oml.datasets.get_dataset(1471)
X, y = dataset.get_data(target=dataset.default_target_attribute)
clf = neighbors.KNeighborsClassifier(n_neighbors=1)
clf.fit(X, y)
```

Out[34]: `KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
metric_params=None, n_jobs=1, n_neighbors=1, p=2,
weights='uniform')`

You can also ask which features are categorical to do your own encoding

```
In [35]: from sklearn import preprocessing
dataset = oml.datasets.get_dataset(10)
X, y, categorical = dataset.get_data(
    target=dataset.default_target_attribute,
    return_categorical_indicator=True)
print("Categorical features: %s" % categorical)
enc = preprocessing.OneHotEncoder(categorical_features=categorical)
X = enc.fit_transform(X)
clf.fit(X, y)
```

Categorical features: [True, True, True, True, True, True, True, True, False, False, True, True,

```
Out[35]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
    metric_params=None, n_jobs=1, n_neighbors=1, p=2,
    weights='uniform')
```

3 Tasks

To run benchmarks consistently (also across studies and tools), OpenML offers Tasks, which include specific train-test splits and other information to define a scientific task. Tasks are typically created via the website by the dataset provider.

3.1 Listing tasks

```
In [36]: task_list = oml.tasks.list_tasks(size=5000) # Get first 5000 tasks
```

```
mytasks = pd.DataFrame(task_list).transpose()
print("First 5 of %s tasks:" % len(mytasks))
print(mytasks.columns)
```

First 5 of 5000 tasks:

```
Index(['MajorityClassSize', 'MaxNominalAttDistinctValues', 'MinorityClassSize',
      'NumberOfClasses', 'NumberOfFeatures', 'NumberOfInstances',
      'NumberOfInstancesWithMissingValues', 'NumberOfMissingValues',
      'NumberOfNumericFeatures', 'NumberOfSymbolicFeatures', 'cost_matrix',
      'did', 'estimation_procedure', 'evaluation_measures', 'name',
      'number_samples', 'quality_measure', 'source_data',
      'source_data_labeled', 'status', 'target_feature',
      'target_feature_event', 'target_feature_left', 'target_feature_right',
      'target_value', 'task_type', 'tid', 'time_limit', 'ttid'],
      dtype='object')
```

```
In [37]: mytasks = mytasks[['tid', 'did', 'name', 'task_type', 'estimation_procedure', 'evaluation_measures']]
print(mytasks.head())
```

	tid	did	name	task_type	estimation_procedure \
1	1	1	anneal	Supervised Classification	10-fold Crossvalidation
2	2	2	anneal	Supervised Classification	10-fold Crossvalidation
3	3	3	kr-vs-kp	Supervised Classification	10-fold Crossvalidation
4	4	4	labor	Supervised Classification	10-fold Crossvalidation
5	5	5	arrhythmia	Supervised Classification	10-fold Crossvalidation

	evaluation_measures
1	predictive_accuracy
2	predictive_accuracy
3	predictive_accuracy
4	predictive_accuracy
5	predictive_accuracy

Search for the tasks you need

```
In [38]: print(mytasks.query('name=="eeg-eye-state"'))
```

	tid	did	name	task_type \
9983	9983	1471	eeg-eye-state	Supervised Classification
14951	14951	1471	eeg-eye-state	Supervised Classification

	estimation_procedure	evaluation_measures
9983	10-fold Crossvalidation	predictive_accuracy
14951	10-fold Crossvalidation	NaN

3.2 Download tasks

```
In [39]: task = oml.tasks.get_task(14951)
         pprint(vars(task))
```

```
{'class_labels': ['1', '2'],
 'cost_matrix': None,
 'dataset_id': 1471,
 'estimation_parameters': {'number_folds': '10',
                           'number_repeats': '1',
                           'percentage': '',
                           'stratified_sampling': 'true'},
 'estimation_procedure': {'data_splits_url': 'http://www.openml.org/api_splits/get/14951/Task_14',
                           'parameters': {'number_folds': '10',
                                           'number_repeats': '1',
                                           'percentage': '',
                                           'stratified_sampling': 'true'},
                           'type': 'crossvalidation'},
 'evaluation_measure': None,
 'target_name': 'Class',
 'task_id': 14951,
```



```
'task_type': 'Supervised Classification'}
```

4 Runs: Train models on tasks

We can run (many) scikit-learn algorithms on (many) OpenML tasks.

```
In [40]: task = oml.tasks.get_task(14951)
         clf = neighbors.KNeighborsClassifier(n_neighbors=1)
         run = oml.runs.run_task(task, clf)
         run.model
```

```
Out[40]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=1, n_neighbors=1, p=2,
                             weights='uniform')
```

Share the run on the OpenML server

```
In [41]: myrun = run.publish()
         print("Uploaded to http://www.openml.org/r/" + str(myrun.run_id))
```

Uploaded to <http://www.openml.org/r/1846520>

4.1 All together

Train any model on any OpenML dataset and upload to OpenML in a few lines of code

```
In [44]: task = oml.tasks.get_task(14951)
         clf = neighbors.KNeighborsClassifier(n_neighbors=5)
         run = oml.runs.run_task(task, clf)
         myrun = run.publish()
         print("Uploaded to http://www.openml.org/r/" + str(myrun.run_id))
```

Uploaded to <http://www.openml.org/r/1846522>

4.2 A Challenge

We'll see many machine learning algorithms in this course. Try to build the best possible models on several OpenML tasks, and compare your results with the rest of the class, and learn from them. Some tasks you could try (or browse openml.org):

- EEG eye state: data_id:[1471](#), task_id:[14951](#)
- Volcanoes on Venus: data_id:[1527](#), task_id:[10103](#)
- Walking activity: data_id:[1509](#), task_id: [9945](#), 150k instances
- Covertypes (Satellite): data_id:[150](#), task_id: [218](#). 500k instances
- Higgs (Physics): data_id:[23512](#), task_id:[52950](#). 100k instances, missing values

Easy benchmarking:

```
In [45]: import openml as oml
        from sklearn import neighbors

        for task_id in [14951, 10103, 9945]:
            task = oml.tasks.get_task(task_id)
            data = oml.datasets.get_dataset(task.dataset_id)
            clf = neighbors.KNeighborsClassifier(n_neighbors=5)
            run = oml.runs.run_task(task, clf)
            myrun = run.publish()
            print("kNN on %s: http://www.openml.org/r/%d" % (data.name, myrun.run_id))

kNN on eeg-eye-state: http://www.openml.org/r/1846523
kNN on volcanoes-a1: http://www.openml.org/r/1846524
kNN on walking-activity: http://www.openml.org/r/1846525
```

4.3 Other possibilities

OpenML's Python API is currently still under development. To be added soon:

- Support for uploading pipelines
- Organizing data sets, algorithms, and experiments into studies
- Downloading previous experiments, evaluations and models
- Uploading new datasets to OpenML
- Filters for listings (e.g. filter by author, tags, other properties)

All of this is already possible with the R and Java API.